

# Stellar Core Data Flow

 @graydon\_pub, Developer @stellarorg



**Show me your flowcharts and conceal your  
tables, and I shall continue to be mystified.  
Show me your tables, and I won't usually need  
your flowcharts; they'll be obvious.**

Fred Brooks

# Talk Overview

This is a talk about the *data* that stellar-core deals with.  
It does not discuss SCP, Horizon, or applications built on Stellar.  
It does not discuss cryptography, finance or trust.

**It is to help you figure out  
*what* is stored and transmitted *where*.**

# Talk overview

1. Review: replicated state machines
2. Data types
3. Data formats
4. Places data lives
5. Movement of data
6. Bonus: external access to data

# **1. Review: replicated state machines**



stellar-core is a  
***replicated state machine***

# State machine

Pure function of current state + input

$$F(\text{State}_n, \text{Input}_{n+1}) \dashrightarrow (\text{State}_{n+1}, \text{Output}_{n+1})$$

Deterministic

Same state + input always makes same next-state + output

Can replay any step, given state + input



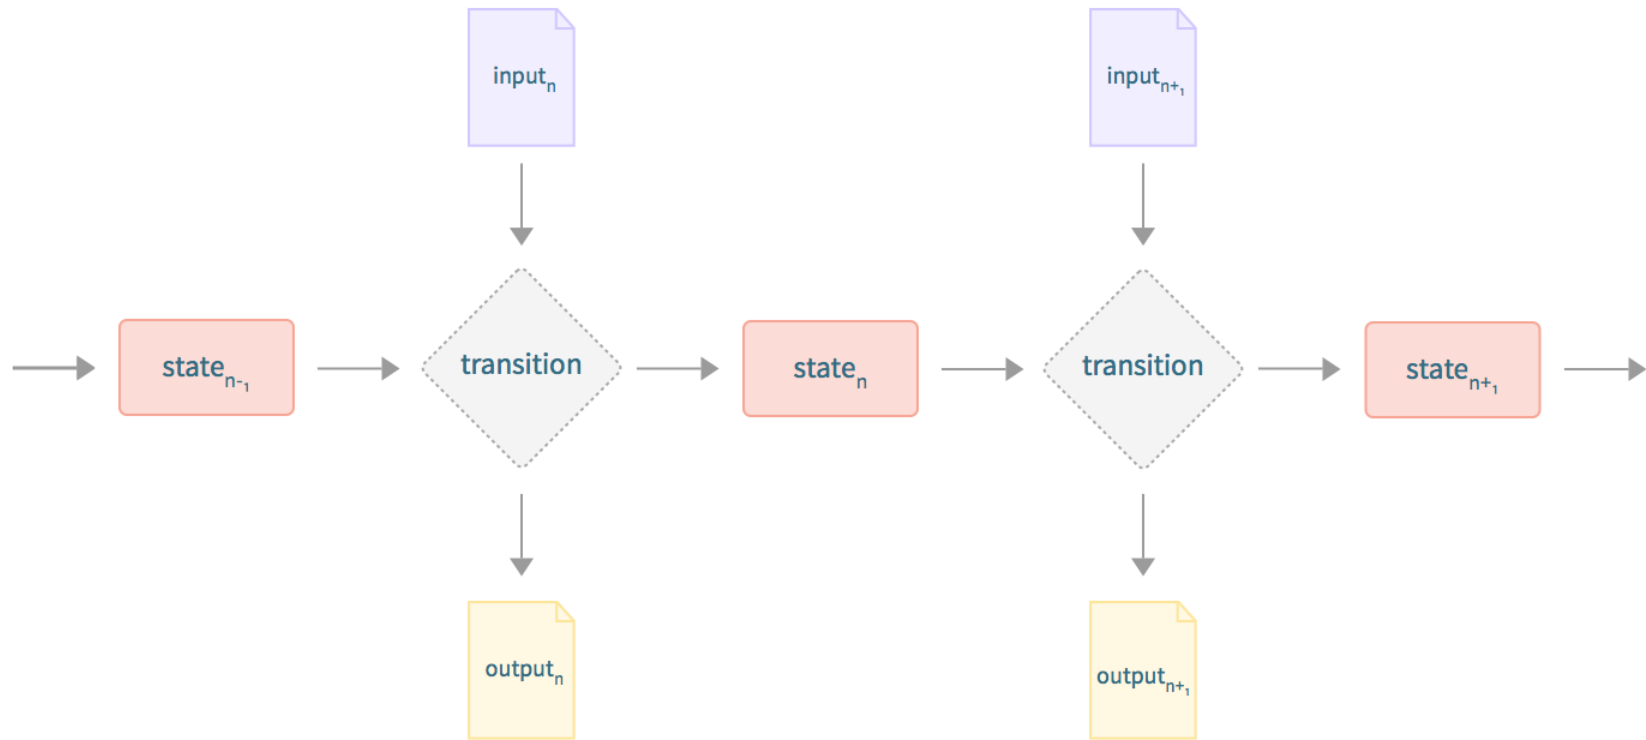
# State machine

We will not discuss the function  $F$  much.

Suffice to say it's "applying transactions."

The other 3 parts are *data*, which we'll talk about:

1. State
2. Input
3. Output



# Replicas

**Recall:** stellar-core is intended as a *replicated* state machine

**Meaning:** keep multiple copies of state machine *and its data*



- On different physical computers
- Run at *same* time, in lock-step
- Run *same* function on *same* "input + state" data
- Produce *same* "output + next state" data

# **Replication is for reliability, decentralization**


"lots of copies keeps stuff safe"

# Replicas are coordinated by a *consensus algorithm*

ensures same current state and input on all replicas



consensus  
algorithm



# **stellar-core uses SCP for replica consensus**

this talk is not about SCP

# In the stellar-core state machine:

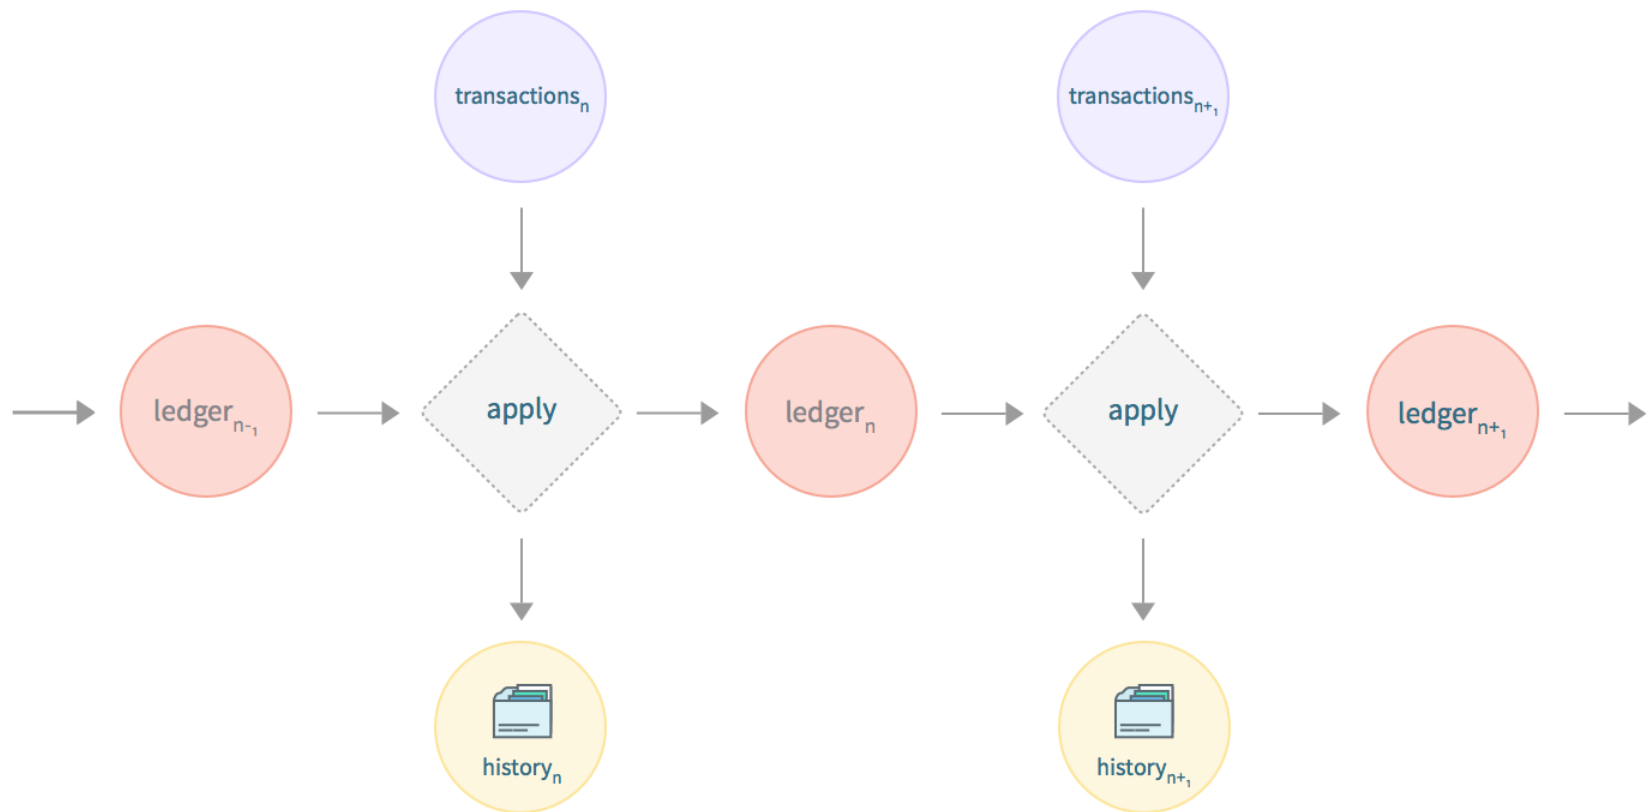
State = Ledger

Input = Transactions

Output = History

$$F(\text{Ledger}_n, \text{Transactions}_{n+1}) \dashrightarrow (\text{Ledger}_{n+1}, \text{History}_{n+1})$$





# Every stellar-core peer follows this cycle

(endless loop, every 5 seconds)



1. Acquire consensus on state and input
2. Apply input to state
3. Emit output, advance to new state

# Every stellar-core peer follows this cycle

(endless loop, every 5 seconds)



1. Acquire consensus on **ledger** and **transactions**
2. Apply **transactions** to **ledger**
3. Emit **history**, advance to new **ledger**

## **2. Data Types**

# Recall: data of stellar-core state machine

1. **Ledgers (state)**
2. **Transactions (input)**
3. **History (output)**

# Ledger

**Recall:** this is the *state* data

Description of how-things-are at the present moment

Set of 3 kinds of *entries*:

- Accounts
- Trustlines
- Offers

# Transactions

**Recall:** this is the *input* data

Is truly *data*: encoded descriptions of actions-to-perform

Handful of possible actions on ledger entries:

- Create/modify/delete entry
- Transfer amount between entries
- Miscellaneous others (inflation, set options, etc.)

# History

**Recall:** this is the *output* data

*Log* of changes during each state-transition:

- Transaction set that was used as *input*
- Success or failure of each transaction, and its effects
- Compact description of *next state*



# **3. Data Formats**

# Data in stellar-core takes 2 forms:

**1. XDR**

**2. SQL**

plus a few auxiliary TOML and JSON files

# XDR

External Data Representation

Generic binary serialization format<sup>1</sup>

Internet standard<sup>2</sup>

Driven by plain-text schemas

<sup>1</sup> Like ASN.1, Protocol Buffers, Thrift, Avro

<sup>2</sup> RFC 4506 / STD 67

# XDR

Structured Query Language

Generic relational database access format

International standard<sup>3</sup>

Implies: **stellar-core always paired with a database**<sup>4</sup>

<sup>3</sup> Like ASN.1, Protocol Buffers, Thrift, Avro

<sup>4</sup> RFC 4506 / STD 67

# Uses of XDR

All 3 kinds of data in stellar-core are expressed in XDR:

- Transactions (input) received in XDR
- Ledger (state) stored on disk in XDR
- History (output) emitted in XDR

Plus all SCP and P2P network messages

# Uses of SQL

*Mostly*<sup>5</sup> just the ledger (state)

*Mostly*<sup>6</sup> just read / written while applying transactions

<sup>5</sup> Some history (output) is also buffered there, on the way out

<sup>6</sup> Consensus does some reading in order to validate potential input

**Wait, isn't the ledger in XDR?**  
**Yes: the ledger is stored *twice***  
**In XDR *and* SQL, simultaneously**

for two good reasons—we'll get to them

## **4. Places data lives**



# Stellar-core deals with data in 4 places

1. XDR in flight (between replicas)
2. SQL tables in a relational database
3. XDR files on local disk
4. XDR files in a "history archive"

# XDR in flight

Peer-to-peer network *between replicas*

Messages flood to all peers

Mainly transactions & SCP messages

Held in memory until consensus

# SQL tables in a relational database

Consulted during consensus

*Modified* during state-machine transition

Modified *atomically*:  $Ledger_n \dashrightarrow Ledger_{n+1}$

Random-access, fine-grained

Fast: hundreds to thousands of updates per second

# XDR files on local disk

So-called "buckets"

Store the ledger in *canonical form*

*Duplicate* of data stored in SQL tables

Needed for 2 operations<sup>7</sup>:

- Efficient, incremental *cryptographic hashing*
- Efficient, incremental *storage and transmission of differences*

<sup>7</sup> See <https://github.com/stellar/stellar-core/blob/master/src/bucket/BucketList.h>

# XDR files in a "history archive"

Long-term, flat-file, mostly cold storage

User-defined backends<sup>8</sup>

Stores *checkpoints*: XDR buckets and XDR history logs

*Mostly*<sup>9</sup> write-once, read-many

Used by peers to catch up to one another

<sup>8</sup> Typically AWS S3, Google Cloud Storage, Azure Blob Storage, SCP/SFTP, etc.

<sup>9</sup> A single JSON file is rewritten to point to the "most recent" checkpoint

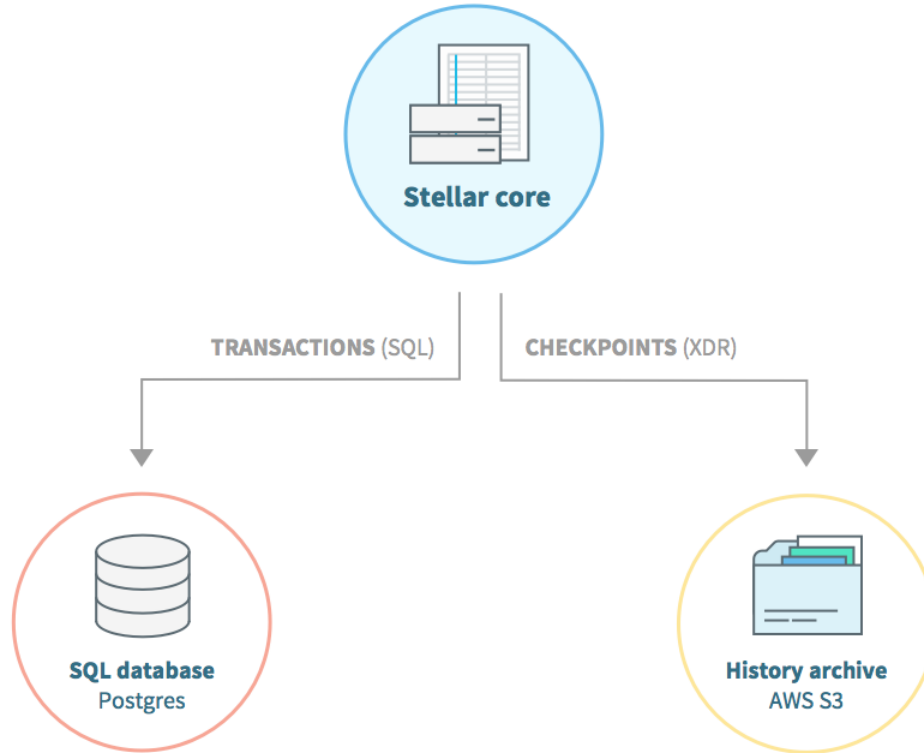
# Reiteration in case this was not clear

A stellar-core node **usually requires** two other storage facilities:

- A relational database<sup>10</sup>
- One or more history archives<sup>11</sup>

<sup>10</sup> SQLite is bundled and may be sufficient for small networks; PostgreSQL is recommended.

<sup>11</sup> At least configuring an archive to *read* from; *writing* to an archive is optional, but recommended.



## **5. Movement of data**



# Data moves in 5 interesting flows

1. History archives  $\rightarrow$  Peers = **"Catchup"**
2. External clients  $\rightarrow$  Peers = **"Submission"**
3. Peers  $\rightarrow$  Peers = **"Flooding"**
4. Peers  $\rightarrow$  Databases and local files = **"Applying"**
5. Peers  $\rightarrow$  History archives = **"Publishing"**

# Catchup

Happens when a peer is new or out of sync

Downloads<sup>12</sup> XDR history files from history archive

One of two operator-chosen modes, either:

- replays state-transitions in order, or
- snaps to most recent state<sup>13</sup>

<sup>12</sup> Archive-specific, configured by user. Usually HTTP GET or similar.

<sup>13</sup> This mode only downloads differences, one of the two reasons for duplicating the ledger in buckets.

# Submission

Happens when an external client has new transaction

Contacts peer through HTTP (likely via Horizon)

Sends XDR representation of transaction

Receives status code indicating "rejected" or "pending"

# Flooding

Happens continuously

Peers hold long-lived TCP connections to one another

All messages are XDR, repeated to all peers

- Transactions: flood as they're submitted

- SCP messages: a burst of activity every 5 seconds

# Applying

Happens when SCP decides on consensus state + input

About every 5 seconds

Transactions in memory applied to ledger in SQL database

Duplicate copies of changed ledger entries put in XDR buckets<sup>14</sup>

Transactions and results written to accumulating XDR checkpoint

<sup>14</sup> Cryptographic hash of ledger is efficiently calculated here: the other reason for duplicating the ledger in buckets.

# Publishing

Happens every 64 ledgers

About every 5 minutes

Uploads<sup>15</sup> accumulated checkpoint to history archive

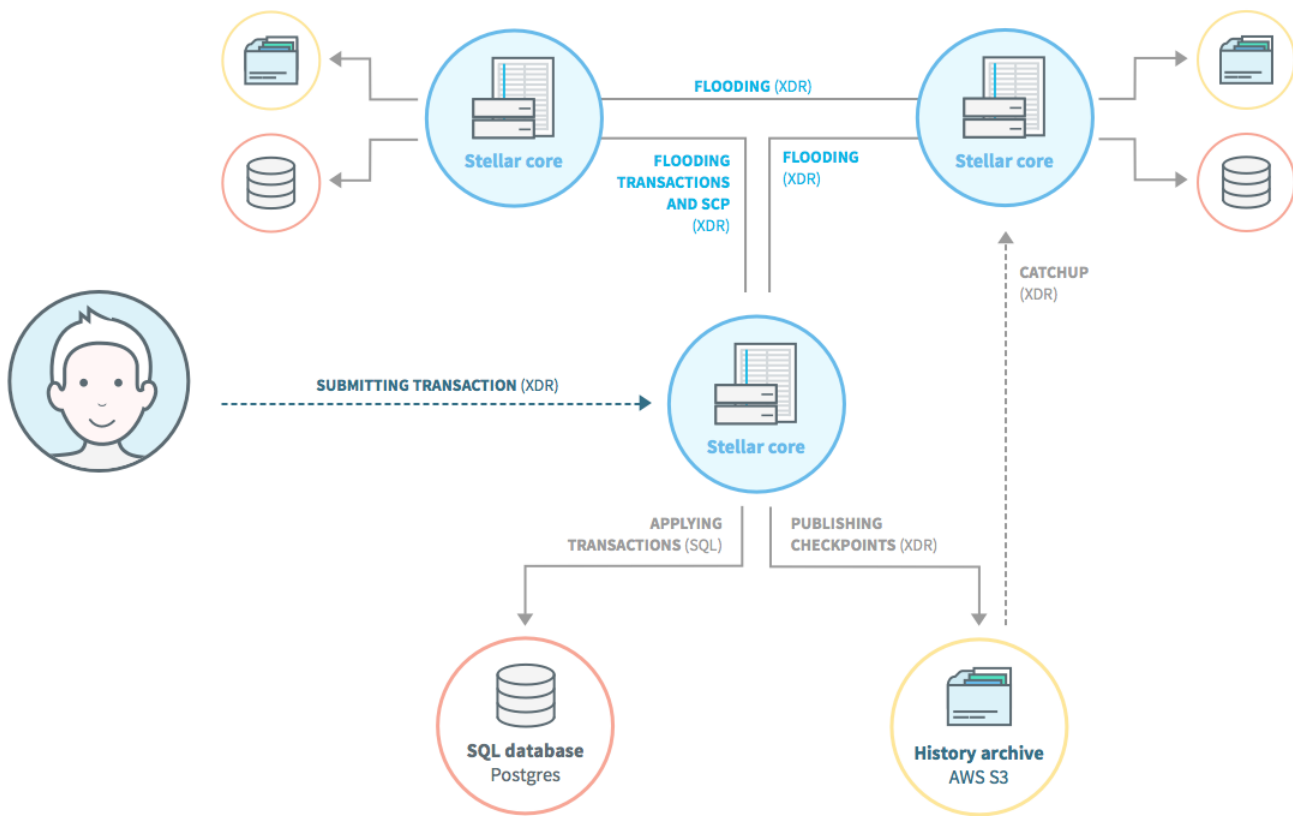
Includes 64 state-transitions worth of history, compressed

All transactions, results, and any new buckets<sup>16</sup>

<sup>14</sup> Archive-specific, configured by user. Usually HTTP PUT or similar.

<sup>15</sup> Only sends buckets differing from previous checkpoints.

**Most-complicated  
diagram time!**





## **6. Bonus: external access to data**

**It may seem a little odd that basic functions like "catchup" go through history archives.**

# History archives serve several roles

Ensuring reliable backups are made

Minimizing risks of single-node failure

Controlling storage costs for largest data set (history)

Isolating catchup I/O load away from P2P flooding

**Providing very simple external access to data**

# A moment about that last point

Stellar is intended as a broadly interoperable system

Simplicity, transparency, standardization are key

Want there to be zero barriers to "getting the data"

- Even if consensus network is offline

- Even if stuck behind a firewall

- Even if polling via shell scripts and duct tape

# Benefits of flat files

You do not need to "talk to" stellar-core to get data

Command-line tools can download from history archives

`curl/wget` usually fine

Reading/interpreting involves only gzip, JSON and XDR

stellar-core will dump an XDR file as plain text, offline

Decoding XDR is pretty straightforward anyways

# Go forth and experiment!

XDR schemas are public<sup>17</sup>

Archives are just directories full of XDR files

If you want to see the transactions in ledger 0x3127:



*/transactions/00/00/31/transactions-00003127.xdr.gz*

<sup>17</sup> See <https://github.com/stellar/stellar-core/tree/master/src/xdr>

# Fini

